CONCENTRATION & EXIT RISK ASSESSMENT

AWS exit risk review: demo organisation

Synthetic data for demonstration. Not a real organisation.

SCOPE	ORGANIZATIONAL UNIT	ACCOUNTS IN SCOPE	SPEND PERIOD
name contains 'prod'	Root	40 of 49	Oct 2025 to Sep 2026 (12 months, amortised)
SCANNED	**MANAGEMENT ACCOUNT**	**RESOURCE INVENTORY**	**EXIT TARGET ASSUMED**
Oct 22, 2026, 2:47 PM	111111111111	Included (deep scan)	Microsoft Azure

1 Executive summary

BOTTOM LINE

Exit risk is Moderate. Most workloads are portable, but the size of the estate makes an exit a large programme. Cost and time are not yet estimated for your organisation: answer the delivery questions under Exit model.

Scores 39/100 by spend and 28/100 by number of accounts. Risk is concentrated: 6 accounts (36% of spend), mostly on Redshift, DynamoDB and Lambda, carry most of it, while 34 of 40 accounts (85%) are broadly portable.

Cost and time are provisional. They use default answers, not your organisation's. Answer the delivery questions (CI/CD, Terraform, team size, day rate) to estimate them for you.

WHAT WE FOUND

- **27% of spend ($2.7m a year)** runs on 19 bespoke, AWS-only services that must be re-engineered: Redshift $746k a year (1 account), DynamoDB $530k (6) and Lambda $391k (10); 16 others $1.0m.
- **34 of 40 accounts are portable** (Low or Moderate risk). The 6 High or Critical accounts hold 36% of spend.
- **3 accounts set the pace:** data-platform-prod, cards-processing-prod and contact-centre-prod carry 57% of the spend at risk (spend x risk).
- **Spend grew 42%** from Oct 2025 to Sep 2026; High/Critical spend grew 38%, so lock-in is growing more slowly than the estate. Biggest monthly rises: EC2 +$62k, RDS / Aurora +$38k and Redshift +$23k.
- **The exit does not fit the 24-month deadline:** it takes roughly 2.5 to 5 years on default delivery answers. Start earlier, add capacity, or agree a longer transition.
- **Commitments outlast the exit:** up to $100k of Savings Plans and reservations would still be owed when it finishes, a write-off; the last commitment ends Sep 2029.
- **90% of the spend of 5 critical accounts** runs in Europe (London), so one regional event takes down most of it at once. Channels Card payments, Core banking, Mobile app, Payments and Shared platform share that region: no anti-affinity between them.

WHAT WE RECOMMEND

1. **Answer the delivery questions under Exit model** (CI/CD, Terraform, team size, day rate). Until then the cost and time figures use default answers, not your organisation's.
2. **Plan the hardest accounts first.** Write and test exit plans for data-platform-prod, cards-processing-prod and contact-centre-prod; they drive most of the cost and the timeline.
3. **Make portability a design criterion** for new work on Redshift, DynamoDB and Lambda: prefer the open-standard alternative in Appendix A unless the business case says otherwise.
4. **Pre-build a stressed-exit fallback** for the 24% of critical spend on AWS-only services: a stand-in process, a manual workaround or a warm copy elsewhere, tested before it is needed.
5. **Match commitment terms to the exit plan:** $318k a year of commitments runs past the fastest exit. Buy new Savings Plans and reservations on terms that end inside the exit window.

Confidence: Low Synthetic demo data. Cost and time use default delivery answers, not yours. 1 service not in the catalog, scored as Moderate.

To raise confidence

1. Scan your own organisation.
2. Answer the delivery questions under Exit model (about ten minutes).
3. Add the missing services to the service catalog (Appendix A lists them).

KEY FIGURES

OU exit-risk score

By spend	By no. of accounts
39 /100	28 /100
Moderate	Moderate

Risk is concentrated in a few large accounts: most accounts are easier to move than the spend view suggests.



Annual spend in scope
$9.9m
$828k / month, amortised

Spend on High/Critical services
27%
$2.7m a year

Accounts in scope
40
40 with spend · 9 filtered out

High/Critical accounts
6 of 40
36% of spend · 34 portable (Low/Moderate)

Bespoke (AWS-only) services
19
largest: Redshift, DynamoDB, Lambda

Spend growth
+41.5%
Oct 2025 → Sep 2026 · High/Critical +38.5%

Exit cost (one-off)
$7m to $16m PROVISIONAL
70 to 180% of a year's AWS spend

Exit window (time to leave AWS)
2.5 to 5 years PROVISIONAL
bound by team capacity

Engineering effort
25 to 60 person-years PROVISIONAL
team of 15 full-time people

Commitments written off
up to $100k PROVISIONAL
still owed after the exit; whole organisation, even when filtered

2 Where the exit risk sits

Monthly spend by exit-risk tier

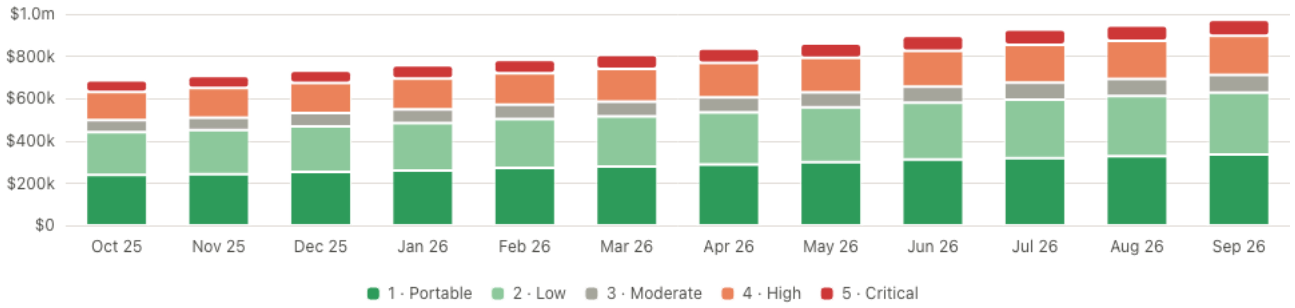
Share of in-scope spend by how hard it is to move.



Tier	Per month	Share	Services	Accounts
1 · Portable	$288k	34.8%	12	35
2 · Low	$245k	29.6%	18	39
3 · Moderate	$70k	8.5%	11	7
4 · High	$159k	19.2%	16	14
5 · Critical	$66k	7.9%	3	7

Spend trend by tier

Is lock-in growing? Monthly amortised spend.

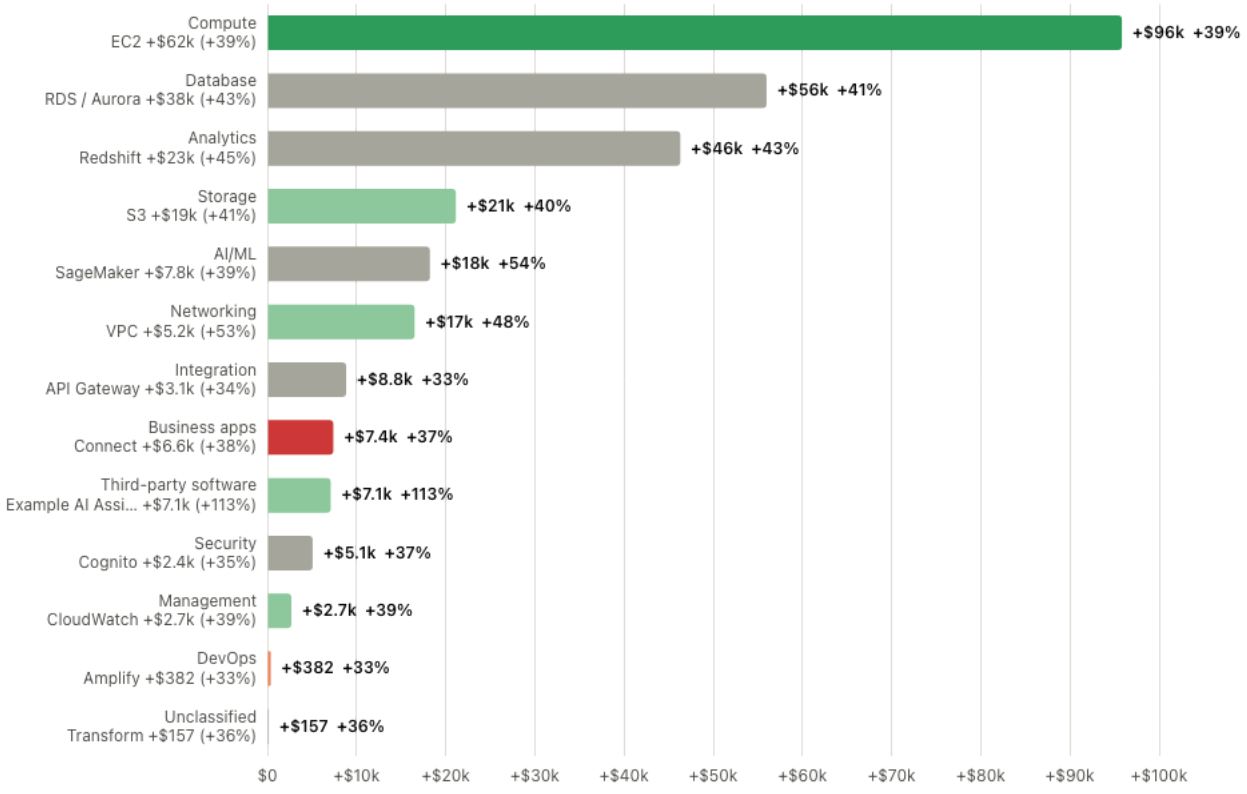


SPEND GROWTH BY PRODUCT TYPE

Oct 2025 to Sep 2026: total spend +41.5% (+46.0% a year); spend on High/Critical services +38.5% (+42.6% a year). Lock-in is growing more slowly than the estate.

Change in monthly spend

Change in monthly spend from Oct 2025 to Sep 2026, by product type, in dollars and percent. Under each type, the service that moved most. Colour = how hard the product type is to exit, from its services' tiers weighted by spend: green easy to exit, grey moderate, red hard.



Monthly spend by product type

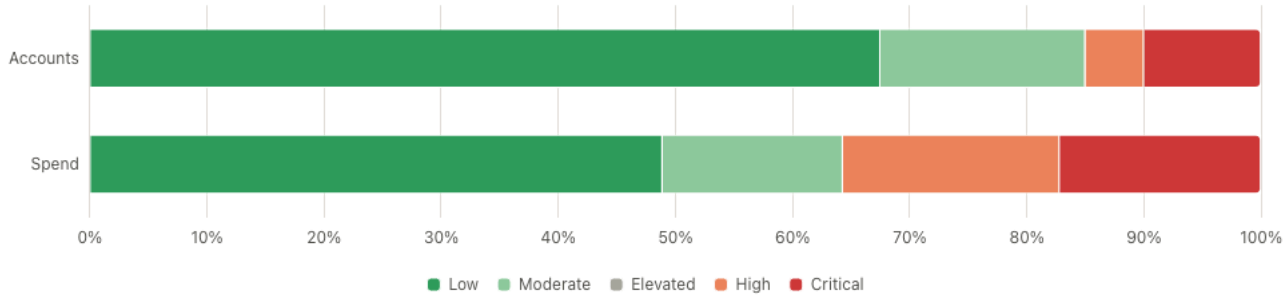
Product type	Oct 2025	Sep 2026	Change	Per year	Lock-in
Compute EC2, EC2 Other (EBS, NAT, transfer), Lambda, ECS +2 more	$247k	$343k	+38.8%	+43.0%	11%
Database RDS / Aurora, DynamoDB, Neptune, ElastiCache +1 more	$137k	$193k	+40.8%	+45.2%	28%
Analytics Redshift, Glue, EMR, Kinesis Streams +9 more	$108k	$154k	+43.1%	+47.8%	71%
Storage S3, FSx, Backup	$54k	$75k	+39.5%	+43.8%	0%
AI/ML SageMaker, Bedrock, Transcribe, Lex +2 more	$34k	$52k	+54.2%	+60.4%	10%
Networking VPC, Direct Connect, CloudFront, Network Firewall +3 more	$34k	$51k	+47.9%	+53.3%	0%
Integration API Gateway, Step Functions, AppSync, SQS +4 more	$27k	$36k	+32.8%	+36.3%	45%
Business apps Connect, Pinpoint	$20k	$27k	+37.2%	+41.2%	100%
Security Cognito, KMS, Directory Service, ACM / Private CA +2 more	$14k	$19k	+36.8%	+40.7%	49%
Third-party software Marketplace (3rd party)	$6.3k	$13k	+112.5%	+127.5%	0%
Management CloudWatch	$7.0k	$9.7k	+38.6%	+42.8%	0%
DevOps Amplify	$1.1k	$1.5k	+33.4%	+36.9%	100%
Unclassified Transform	$435	$592	+36.0%	+39.9%	0%
Total	$689k	$974k	**+41.5%**	+46.0%	

Per year = the first-to-last-month change annualised. Amortised spend: Savings Plans and RIs are spread over the months they cover.

3 Accounts: how many are at risk, and how big

Accounts by exit-risk rating

Share of accounts (top) against share of spend (bottom). Few accounts holding a lot of spend is a concentration to manage, not an estate-wide problem.



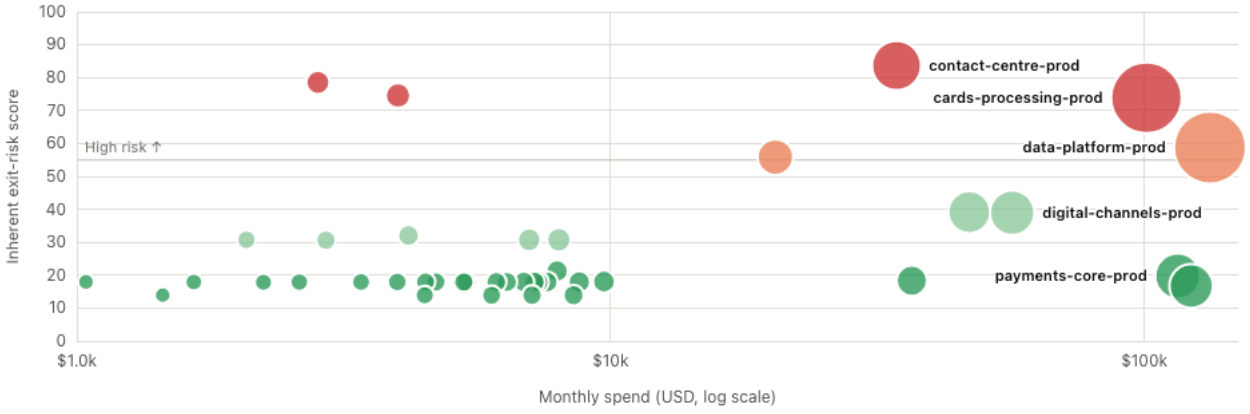
Portable vs higher-risk accounts

Rating	Accounts	Share	Per month	Spend share
Low	27	67.5%	$405k	48.9%
Moderate	7	17.5%	$128k	15.4%
Elevated	0	0.0%	$0	0.0%
High	2	5.0%	$153k	18.5%
Critical	4	10.0%	$142k	17.2%
Portable (Low + Moderate)	**34**	85.0%		

OU score weighting: By spend. Each account counts in proportion to its spend, so a small high-risk account adds little. It sets bubble sizes and contributions; change it in the bar at the top.

Inherent risk vs spend

Each bubble is an account. Up = harder to move; right = bigger; bubble size = points it contributes to the OU score. Top-right accounts drive the exit.



Account ranking by criticality: spend × risk

#	Account	Per month	Risk	Spend at risk / month	Main lock-in
1	**data-platform-prod** 561017993421 · Root / Workloads / Production · Owner: Emma Botha	$133k	High · 59 63% High/Crit spend	$78k	Redshift Glue QuickSight
2	**cards-processing-prod** Critical 966773840736 · Card payments · Root / Workloads / Production · Owner: Priya Shah	$101k	Critical · 74 77% High/Crit spend	$75k	DynamoDB Lambda Kines
3	**contact-centre-prod** 125285552493 · Contact centre · Root / Workloads / Production · Owner: Lerato Mokoena	$34k	Critical · 84 86% High/Crit spend	$29k	Connect Lex Lambda
4	**payments-core-prod** Critical 534439589175 · Payments · Root / Workloads / Production · Owner: Priya Shah	$116k	Low · 20 3% High/Crit spend	$23k	DynamoDB Lambda
5	**digital-channels-prod** Critical 978681853473 · Mobile app · Root / Workloads / Production · Owner: Lerato Mokoena	$57k	Moderate · 39 25% High/Crit spend	$22k	Cognito AppSync Pinpoin
6	**core-banking-prod** Critical 813958344723 · Core banking · Root / Workloads / Production · Owner: Thabo Dlamini	$122k	Low · 17 0% High/Crit spend	$21k	none
7	**ml-platform-prod** 217230843944 · Root / Workloads / Production · Owner: Emma Botha	$47k	Moderate · 39 3% High/Crit spend	$18k	Bedrock AgentCore
8	**fraud-analytics-prod** 466806226017 · Root / Workloads / Production · Owner: Emma Botha	$20k	High · 56 28% High/Crit spend	$11k	DynamoDB Timestream L
9	**shared-network-prod** 263737281768 · Root / Infrastructure	$37k	Low · 18 0% High/Crit spend	$6.7k	none
10	**search-prod** 899894212661 · Root / Workloads / Production	$4.0k	Critical · 75 70% High/Crit spend	$3.0k	DynamoDB Lambda
11	**mobile-backend-prod** 434423300025 · Root / Workloads / Production · Owner: Lerato Mokoena	$8.0k	Moderate · 31 15% High/Crit spend	$2.5k	Lambda
12	**treasury-reporting-prod** 499647302627 · Root / Workloads / Production · Owner: Thabo Dlamini	$2.8k	Critical · 79 92% High/Crit spend	$2.2k	DynamoDB QuickSight La
13	**merchant-portal-prod** 386906790760 · Root / Workloads / Production	$7.0k	Moderate · 31 15% High/Crit spend	$2.2k	Lambda
14	**doc-store-prod** 454519794043 · Root / Workloads / Production	$9.7k	Low · 18 0% High/Crit spend	$1.8k	none
15	**identity-prod** Critical 189183446249 · Shared platform · Root / Infrastructure · Owner: Thabo Dlamini	$7.9k	Low · 21 0% High/Crit spend	$1.7k	none
16	**id-verify-prod** 267387487076 · Root / Workloads / Production	$8.7k	Low · 18 0% High/Crit spend	$1.6k	none
17	**payroll-api-prod** 295250496497 · Root / Workloads / Production	$7.6k	Low · 18 0% High/Crit spend	$1.4k	none
18	**legacy-batch-prod** 107537456074 · Root / Workloads / Production	$4.2k	Moderate · 32 25% High/Crit spend	$1.3k	SWF Data Pipeline
19	**branch-queue-prod** 717640922470 · Root / Workloads / Production	$7.3k	Low · 18 0% High/Crit spend	$1.3k	none
20	**insurance-quotes-prod** 498633168856 · Root / Workloads / Production	$7.2k	Low · 18 0% High/Crit spend	$1.3k	none
21	**disputes-prod** 844781801768 · Root / Workloads / Production	$6.9k	Low · 18 0% High/Crit spend	$1.2k	none
22	**investments-prod** 514740085648 · Root / Workloads / Production	$8.5k	Low · 14 0% High/Crit spend	$1.2k	none
23	**card-controls-prod** 170174047247 · Root / Workloads / Production	$6.4k	Low · 18 0% High/Crit spend	$1.2k	none
24	**pricing-prod** 322790187289 · Root / Workloads / Production	$6.1k	Low · 18 0% High/Crit spend	$1.1k	none
25	**reporting-hub-prod** 244219422241 · Root / Workloads / Production · Owner: Emma Botha	$7.1k	Low · 14 0% High/Crit spend	$997	none

Top 25 of 40 accounts by spend at risk. The JSON export lists every account.

Who owns the risk from the accounts' owner and executive tags

Chris Ndlovu carries the most spend at risk: $109k a month across 4 accounts.

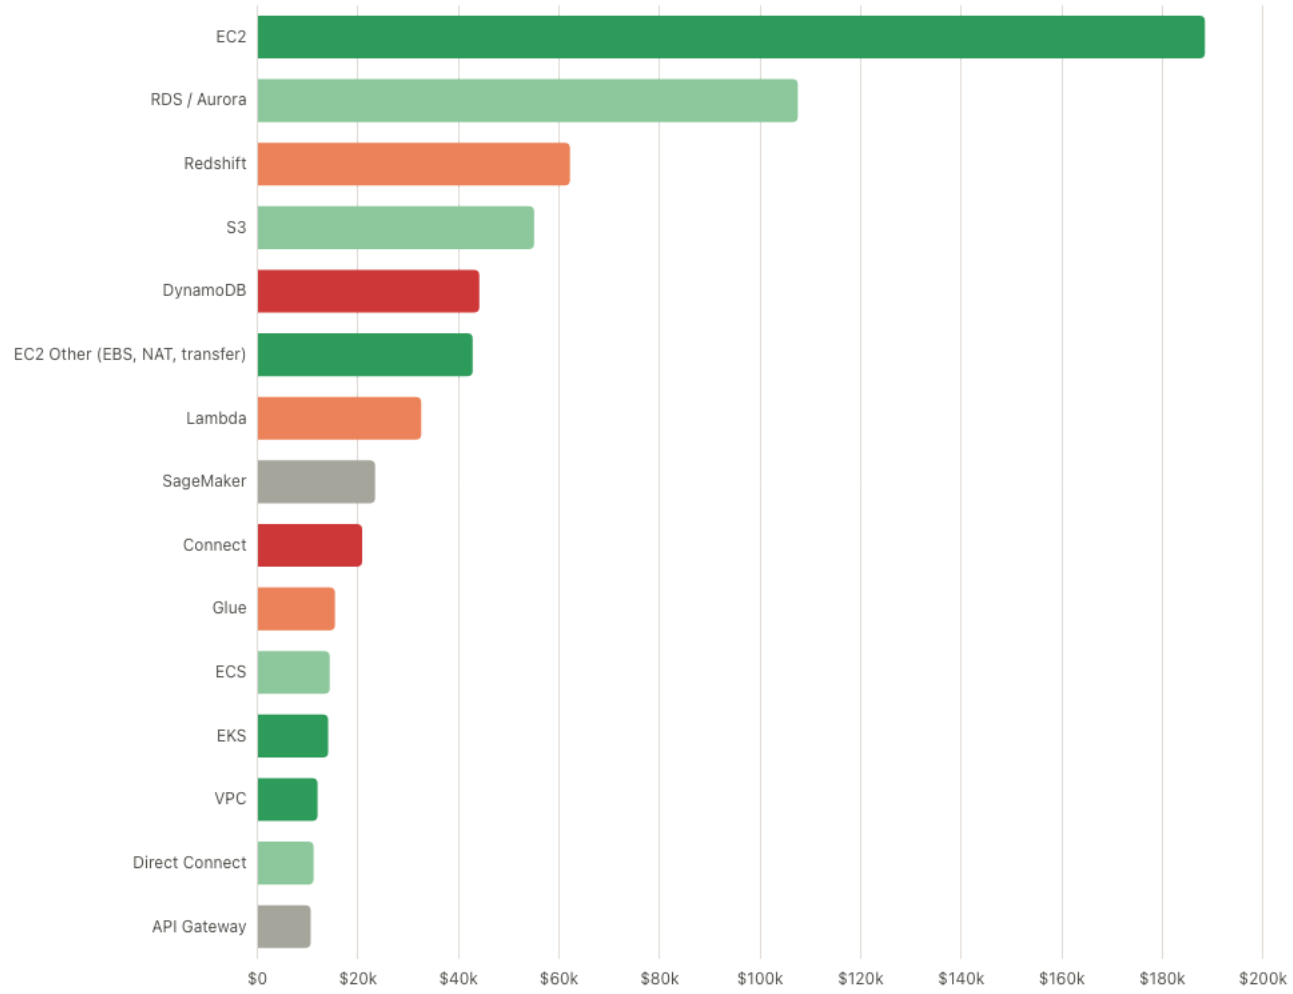
Executive or owner	Accounts	Per month	Spend at risk / month	High/Critical accounts	Account owners
Chris Ndlovu	4	$207k	$109k	2	Emma Botha
Alex Morgan	2	$216k	$97k	1	Priya Shah
Sam Patel	3	$99k	$53k	1	Lerato Mokoena
Jordan Lee	3	$133k	$25k	1	Thabo Dlamini

12 accounts of 40 with spend carry an owner or executive tag. 28 accounts ($172k a month) have none and are left out of this table.

4 Services driving lock-in

Top services by spend

Coloured by exit-risk tier.



High & critical services in use

Service	Tier	Per month	Share	Accounts	Change
Redshift Analytics	4 · High	$62k	7.5%	1	+$23k +45%
DynamoDB Database	5 · Critical	$44k	5.3%	6	+$13k +35%
Lambda Compute	4 · High	$33k	3.9%	10	+$9.8k +35%
Connect Business apps	5 · Critical	$21k	2.5%	1	+$6.6k +38%
Glue Analytics	4 · High	$15k	1.9%	1	+$5.4k +42%
Kinesis Streams Analytics	4 · High	$9.7k	1.2%	2	+$2.8k +34%
Cognito Security	4 · High	$8.0k	1.0%	2	+$2.4k +35%
Step Functions Integration	4 · High	$7.4k	0.9%	2	+$1.8k +27%
QuickSight Analytics	4 · High	$5.9k	0.7%	2	+$2.0k +41%
AppSync Integration	4 · High	$4.3k	0.5%	1	+$1.3k +37%
Lex AI/ML	4 · High	$3.1k	0.4%	1	+$1.1k +44%
Pinpoint Business apps	4 · High	$2.8k	0.3%	1	+$837 +34%
Timestream Database	4 · High	$2.0k	0.2%	1	+$783 +48%
EventBridge Integration	4 · High	$1.7k	0.2%	1	+$510 +36%
Amplify DevOps	4 · High	$1.3k	0.2%	1	+$382 +33%
Bedrock AgentCore AI/ML	4 · High	$1.2k	0.1%	1	+$473 +47%
SWF Integration	5 · Critical	$712	0.1%	1	+$215 +35%
Lake Formation Analytics	4 · High	$640	0.1%	1	+$219 +42%
Data Pipeline Analytics	4 · High	$327	0.0%	1	+$86 +30%

Third-party products billed through AWS

Marketplace software and third-party models on Bedrock. Usually portable: the vendor sells direct or on other clouds, but each contract needs re-papering.

Product	Billed as	Per month	Accounts	Change
Claude Sonnet 4.6	Bedrock	$3.9k	1	+$4.9k
Example Streaming Cloud - Pay As You Go	Marketplace (3rd party)	$3.5k	1	+$1.6k
Example Observability Platform	Marketplace (3rd party)	$2.7k	1	+$2.0k
Example Data Integration Cloud	Marketplace (3rd party)	$1.7k	1	+$564
Example AI Assistant for Enterprise	Marketplace (3rd party)	$1.2k	1	+$3.0k
Claude Haiku 4.5	Bedrock	$1.1k	1	+$963

5 Exit cost & time estimate

Total exit cost (indicative)

$7m to $16m PROVISIONAL

70 to 180% of a year's AWS spend · about 75% engineering, the rest parallel running

Exit window (time to leave AWS)

2.5 to 5 years PROVISIONAL

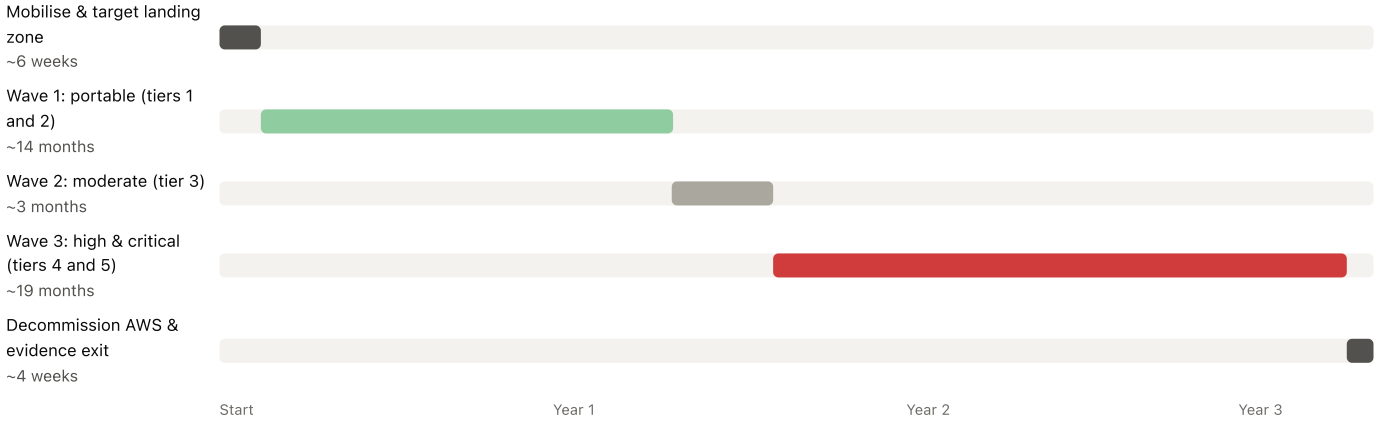
bound by team capacity

Engineering effort

25 to 60 person-years PROVISIONAL

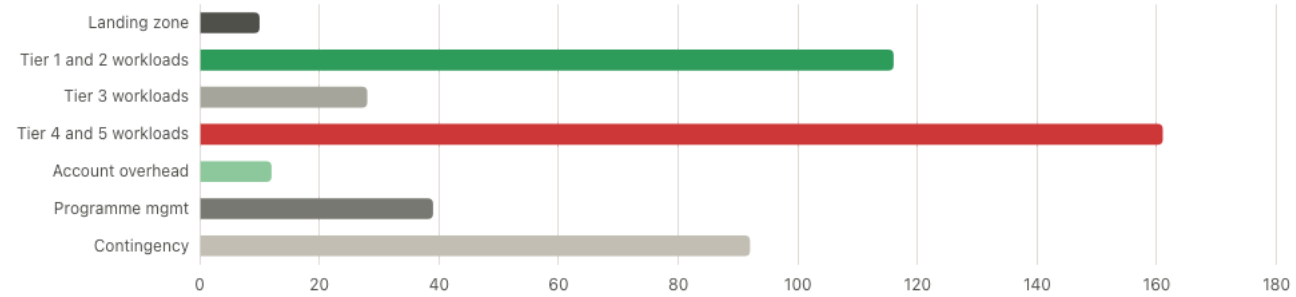
40 material accounts

Indicative timeline



Where the effort goes

Person-months, after the delivery-maturity adjustment.



Your answers and the assumptions behind the estimate

Standard CI/CD pipeline Several, partially standard pipelines	Terraform / IaC Terraform with significant drift	Core cloud team 10 people at 50%	Application engineers 10 FTE
Target Microsoft Azure	Blended day rate 750 USD	Contingency 25%	Parallel running 6 months
Regulated Yes	AI coding agents Not used	Left out Nothing	

- Service effort: catalog engineer-weeks per workload × √(monthly spend / 1,000 USD), clamped to 0.25 to 8.0×; resource counts from the deep scan raise it where higher.
- Services under 25 USD/month are scored but not estimated (immaterial).
- 1 engineer-week per material account for network, IAM and decommissioning.
- Target landing zone: 24 weeks (4 + 12 × √(spend / 100,000 USD), max 24) + 0.5 per material account.
- Delivery factor 1.26 = CI/CD 1.0 × IaC 1.1 × target 1.0 × regulated 1.15.
- Programme management and assurance 12%; contingency 25%.
- Engineers are 80% productive on migration work; duration is the longer of effort ÷ capacity and the critical path of the hardest tier present, plus decommissioning.
- Parallel running: 6 months at 50% of current AWS spend (target ramps up as AWS ramps down).
- Excluded: target-cloud run cost after exit, licence changes, AWS commitments still owed after the exit (EDP, Savings Plans, RIs; shown under Timing), and data-transfer-out fees (since March 2024 AWS credits egress for customers leaving AWS, on request to AWS Support and subject to approval, for a move completed within about 60 days; budget egress for any data moved outside that window).
- Figures are quoted as ranges, widened to round numbers: −25% / +50% on cost and effort, −20% / +40% on duration (early-stage estimation cone). The model's single figures are planning midpoints, not forecasts.

Timing: your deadline and AWS commitments

Notice period	EDP ends	EDP annual commit
24 months	Jun 2028	$8.0m

Exit finishes

Jun 2029 to May 2031 PROVISIONAL

2.5 to 5 years if it starts now

Deadline

Does not fit PROVISIONAL

24-month notice

Commitments

$199k/mo

$4.2m still to run · last ends Sept 2029

EDP

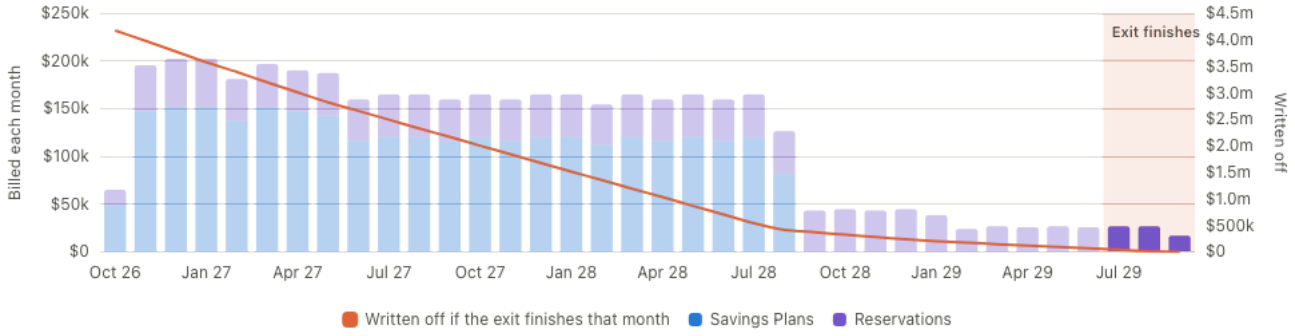
ends Jun 2028

$13.5m of commit left

- The exit does not fit the 24-month notice period: even the fast end runs over, and the slow end is 31 months over.
- Up to $100k of commitments would be written off: still owed after the exit finishes; $27k/mo runs past the fastest finish.
- 3-year terms already outlast the fastest exit; 1-year terms stay safe until Jun 2028.
- The EDP ends Jun 2028, before the exit finishes: renew on a shorter term or without a commit.

Commitments written off, over time

The exit finishes between Jun 2029 and May 2031 (shaded). Bars after the fastest finish are commitments paid for but no longer used: up to $100k written off. The red line is what would be written off if the exit finished at the end of that month. Savings Plans and reservations cover the whole organisation, even though the report is filtered (name contains 'prod').



Savings Plans and reservations

2 Savings Plans and 3 reservations from Cost Explorer, amortised per month. Written off is what would still be owed between the exit finishing and the commitment ending: paid for, no longer used.

Commitment	Term	Ends	Per month	Still to run	Written off
Compute Savings Plan 97% used	3yr No Upfront	Aug 2028	$118k	$2.6m	–
RDS reservation 89% used	3yr All Upfront	Sept 2029	$27k	$927k	$0 to $95k
Redshift reservation 85% used	3yr All Upfront	Jan 2029	$18k	$481k	–
EC2 Instance Savings Plan 93% used	1yr Partial Upfront	May 2027	$31k	$217k	–
ElastiCache reservation 81% used	1yr No Upfront	Feb 2027	$5.2k	$21k	–

6 Resilience: critical accounts, data and a stressed exit

A planned exit is one question. Regulators also ask what happens to critical services if AWS, or one region, fails, and what could move in weeks rather than years. These figures inform; they do not change the score or the estimate.

Critical accounts	Direct Connect today	Carrier cost per new link
5 accounts tagged	2 × 10 Gbps	Not entered

Critical accounts
5
49% of spend · $4.8m a year

Correlated failure
High · 90
90% of critical spend in Europe (London)

No rapid exit
24%
of critical spend is on AWS-only services

Data to move
2.7 PB
13 to 52 days on 1 to 4 new 10 Gbps links

Correlated failure: where the critical estate runs

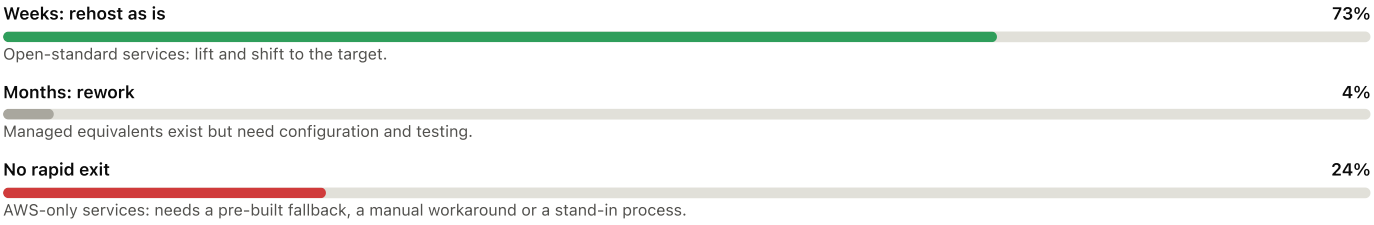
Share of critical spend by region (5 critical accounts). One regional event takes out everything in the top bar at once. A further 3.0% is global services (IAM, Route 53, CloudFront), which share a control plane across regions.



Channel	Runs mainly in	Isolated from other channels?
Card payments	Europe (London) 90%	No
Core banking	Europe (London) 90%	No
Mobile app	Europe (London) 90%	No
Payments	Europe (London) 90%	No
Shared platform	Europe (London) 90%	No

Stressed exit: what could move in a hurry

If AWS stopped serving you with little notice, how fast could critical spend ($4.8m a year) move?



Critical account with no rapid exit	Per month	AWS-only
cards-processing-prod DynamoDB, Lambda, Kinesis Streams	$101k	77%

Data held: what an exit has to move

Average data stored last month, from Cost Explorer storage usage (GB-months); no data is read. 3.3 PB in all: 2.4 PB live, 261 TB archived and 594 TB of snapshots and backups, which are rebuilt at the target rather than moved.

- 2.7 PB (live and archived) has to move; the options for the links that carry it are below.
- Over the internet instead, AWS waives data transfer out for customers leaving AWS if the move finishes within 60 days: that needs about 9 Gbps of internet capacity for two months. Without that credit, internet egress at list price would be roughly $140k to $421k.
- Critical accounts hold 890 TB: the data a stressed exit would need first.
- DynamoDB holds 47 TB; the largest account (cards-processing-prod) holds 37 TB. Exports to S3 are not region-limited, but rebuilding in a second region by import from S3 takes 38 import jobs of up to 1 TB each outside us-east-1, us-west-2 and eu-west-1.
- By region: Europe (London) 3.3 PB.

Service	Live	Archive	Copies
Simple Storage Service	1.6 PB	261 TB	–
Redshift	486 TB	–	121 TB
Relational Database Service	126 TB	–	262 TB
EC2 - Other	209 TB	–	167 TB
DynamoDB	47 TB	–	43 TB

Account	Live	Archive	Copies
data-platform-prod	1.1 PB	261 TB	136 TB
payments-core-prod Critical	216 TB	–	111 TB
core-banking-prod Critical	115 TB	–	143 TB
cards-processing-prod Critical	163 TB	–	34 TB
ml-platform-prod	123 TB	–	–
digital-channels-prod Critical	85 TB	–	24 TB
shared-network-prod	41 TB	–	33 TB
investments-prod	59 TB	–	6.6 TB

Links for the migration: how fast the data can move

Today you have 2 × 10 Gbps of Direct Connect (Equinix LD4 - LD6), costing $3.2k a month in ports and carrying 272 TB out last month. Those carry production traffic and stay as they are. The migration gets its own new 10 Gbps links, dedicated to the copy and kept apart from production, so moving 2.7 PB cannot slow customer-facing systems. Time assumes half the line rate is usable.

Option	New links	Time to copy	AWS ports	Data out	Total
Cheap Lowest cost; the copy takes longest.	1 × 10 Gbps	52 days	$4.9k	$56k	**$61k**
Medium Half the time of Cheap.	2 × 10 Gbps	26 days	$6.6k	$56k	**$63k**
Urgent Fastest: for a deadline or a stressed exit.	4 × 10 Gbps	13 days	$13k	$56k	**$69k**

- AWS list prices for Europe (London) (price list of 2026-10-13): $2.25 per 10 Gbps port-hour, paid for the copy plus 1 month to commission, test and run the final sync; data out over Direct Connect at $0.020 per GB on average, from the regions that hold the data. Data out costs the same whichever option you pick; the links only change how long it takes.
- Not included: your carrier's circuit and the cross-connect at the Direct Connect location. Enter them under Exit model to add them.
- Order early: a new dedicated connection needs a cross-connect at the Direct Connect location and a carrier circuit to your site, which often take weeks to months to deliver.
- AWS's price list also has flat-rate 10 Gbps ports with no per-GB charge, from $10.96 an hour in this region. For a one-off copy of this size, ask AWS whether one costs less than paying per GB.

7 Data quality & scan issues

What this report does *not* cover. Every API error during the scan is recorded here and in the log.

Critical: 0 | Error: 0 | Warning: 2 | Info: 0

Severity	Area	Account	Issue
Warning	Inventory Ap-Southeast-2	data-platform-prod	Eventbridge_rules in ap-southeast-2: Rate exceeded `ThrottlingException` → AWS throttled the request after retries. Re-run, or lower EXITRISK_MAX_WORKERS.
Warning	Catalog	–	'AWS Transform' is not in the service catalog; scored as Moderate (tier 3). $495/month in scope. → Add it to src/exitrisk/data/service_catalog.yaml (see CONTRIBUTING.md).

A Appendix: service exit-risk catalog

Every AWS service is classified into one of five tiers by how much of a workload must change to leave AWS. The tier sets the service's risk score; the recommended path and typical duration are for one workload of that service.

Tier	Score	Meaning
1 · Portable	10	Commodity infrastructure or an open standard. Lift-and-shift; config moves, code does not change.
2 · Low	30	Managed flavour of an open engine or de-facto standard API. Data and config migrate; little code change.
3 · Moderate	50	AWS-specific API with a close equivalent on every major cloud. Code changes at the integration points.
4 · High	75	Proprietary programming or data model. The workload must be redesigned, not just re-pointed.
5 · Critical	95	No equivalent outside AWS. Rebuild, re-platform or replace with a product; often a business change.

Tier 5 · Critical 3 services

Service	Why this tier	Recommended migration path	Equivalents	Typical duration
DynamoDB Database AI agent help: low $44k/mo in scope	Proprietary key-value/document model. Single-table designs, GSIs, conditional writes, TTL and DynamoDB Streams are baked into application code and data layout. No wire-compatible equivalent outside AWS.	Redesign the data access layer. Common targets are Cosmos DB, Firestore/Bigtable, ScyllaDB (Alternator offers a DynamoDB-compatible API) or PostgreSQL JSONB. Replace Streams consumers. Run dual-writes during cutover.	**Azure:** Cosmos DB **GCP:** Firestore / Bigtable / Spanner **Neutral:** ScyllaDB Alternator / Cassandra / PostgreSQL	2-6 months per application
Connect Business apps AI agent help: none $21k/mo in scope	Full contact-centre platform - contact flows, phone numbers, Lex bots, Lambda integrations, recordings and agent desktop. Number porting and agent retraining make it a business transformation.	Select a CCaaS platform (Genesys, NICE, Five9, Dynamics 365 Contact Center), rebuild flows, port numbers, migrate recordings under retention rules, retrain agents.	**Azure:** Dynamics 365 Contact Center **GCP:** Contact Center AI Platform **Neutral:** Genesys Cloud / NICE CXone / Five9	6-12 months
SWF Integration AI agent help: high $712/mo in scope	Legacy proprietary workflow engine with a bespoke decider/worker model; no equivalent.	Rewrite on Temporal (the conceptual successor) or Durable Functions.	**Azure:** Durable Functions **GCP:** Workflows **Neutral:** Temporal	2-4 months

Tier 4 · High 16 services

Service	Why this tier	Recommended migration path	Equivalents	Typical duration
Redshift Analytics AI agent help: low $62k/mo in scope	PostgreSQL-derived but proprietary MPP warehouse - distribution/sort keys, Spectrum external tables, stored procedures, UNLOAD/COPY from S3, Redshift ML and data sharing are all specific. Downstream BI and ELT jobs depend on its dialect.	Choose a warehouse (Snowflake or Databricks are cloud-neutral; BigQuery / Fabric Warehouse are native). Convert DDL and SQL with a translation tool, re-point ELT (dbt makes this far cheaper), reconcile results in parallel for several reporting cycles.	**Azure:** Fabric Warehouse / Synapse / Snowflake **GCP:** BigQuery **Neutral:** Snowflake / Databricks / ClickHouse	3-9 months per warehouse
Lambda Compute AI agent help: high $33k/mo in scope	The function code is portable, but the event model, triggers (S3, SQS, DynamoDB streams, EventBridge), IAM wiring, layers and runtime limits are AWS-specific. Real Lambda estates are mostly glue between other AWS services.	Inventory triggers first. Repackage as containers on Cloud Run / Container Apps / Knative, or rewrite as Azure Functions; replace each event source with the target's equivalent. Migrate with the services it glues together, not alone.	**Azure:** Azure Functions / Container Apps **GCP:** Cloud Run functions **Neutral:** Knative / OpenFaaS on Kubernetes	4-12 weeks per application
Glue Analytics AI agent help: high $15k/mo in scope	Spark under the hood, but jobs use GlueContext/DynamicFrame APIs, the Glue Data Catalog, crawlers and Glue workflows. The catalog is usually shared by Athena, EMR, Redshift Spectrum and Lake Formation, so it is a hub.	Port jobs to plain Spark (Databricks, Dataproc, Synapse Spark), migrate the catalog to Unity / Hive Metastore / Dataplex, and replace crawlers with explicit schemas.	**Azure:** Data Factory + Databricks **GCP:** Dataproc / Dataflow + Dataplex **Neutral:** Spark + Airflow + Hive Metastore	1-4 months
Kinesis Streams Analytics AI agent help: medium $9.7k/mo in scope	Proprietary streaming API (shards, KCL checkpointing in DynamoDB, enhanced fan-out). Producers and consumers must be rewritten.	Move to Kafka (MSK-equivalent, Confluent, Event Hubs Kafka endpoint) or Pub/Sub; rewrite producers/consumers against the Kafka API to avoid the next lock-in.	**Azure:** Event Hubs (Kafka API) **GCP:** Pub/Sub / Managed Kafka **Neutral:** Apache Kafka / Redpanda	1-3 months
Cognito Security AI agent help: medium $8.0k/mo in scope	User pools hold customer identities; password hashes cannot be exported, and triggers / hosted UI are AWS-specific. Customers notice an identity migration.	Move to Entra External ID, Identity Platform, Auth0 or Keycloak using lazy (on-login) migration; plan MFA re-enrolment and app token changes.	**Azure:** Entra External ID **GCP:** Identity Platform **Neutral:** Auth0 / Keycloak / Okta CIC	2-4 months (customer-facing)
Step Functions Integration AI agent help: high $7.4k/mo in scope	Workflows are written in Amazon States Language with direct service integrations to 200+ AWS APIs; nothing executes ASL outside AWS.	Rewrite workflows in Logic Apps / Durable Functions, Google Workflows, or Temporal (cloud-neutral and code-first). Redesign each service integration step.	**Azure:** Logic Apps / Durable Functions **GCP:** Workflows **Neutral:** Temporal / Argo Workflows	1-3 months
QuickSight Analytics AI agent help: medium $5.9k/mo in scope	Proprietary BI tool; dashboards, SPICE datasets and embedding cannot be exported to other BI products.	Rebuild priority dashboards in Power BI, Looker or Superset; retire the long tail. Treat as a business change, not a technical migration.	**Azure:** Power BI **GCP:** Looker **Neutral:** Apache Superset / Tableau	1-3 months

Service	Why this tier	Recommended migration path	Equivalents	Typical duration
AppSync Integration AI agent help: high $4.3k/mo in scope	GraphQL is standard but AppSync resolvers (VTL / APPSYNC_JS), data-source bindings and subscriptions are proprietary.	Re-implement with Apollo Server / Hasura / Azure API Management GraphQL; rewrite resolvers.	**Azure:** API Management (GraphQL) / Hasura **GCP:** Apollo on Cloud Run **Neutral:** Apollo / Hasura	1-3 months
Lex AI/ML AI agent help: medium $3.1k/mo in scope	Proprietary bot definitions, intents and Connect integration.	Rebuild bots in Copilot Studio / Dialogflow CX or an LLM-based assistant.	**Azure:** Copilot Studio **GCP:** Dialogflow CX **Neutral:** Rasa / LLM agents	1-3 months
Pinpoint Business apps AI agent help: medium $2.8k/mo in scope	Proprietary campaign, journey and segment definitions; sender IDs and registered numbers are tied to AWS. AWS ends support for Pinpoint's engagement features (campaigns, journeys, segments, analytics) on 30 October 2026; SMS, voice and push continue as AWS End User Messaging.	Move to Braze / Twilio / Azure Communication Services, or within AWS to Connect outbound campaigns; re-register sender IDs and short codes.	**Azure:** Azure Communication Services **GCP:** Firebase Cloud Messaging + partner **Neutral:** Braze / Twilio	1-3 months
Timestream Database AI agent help: medium $2.0k/mo in scope	Proprietary time-series engine and SQL dialect (Timestream for InfluxDB is the portable variant).	Move to InfluxDB, TimescaleDB, Azure Data Explorer or Bigtable; rewrite queries and ingestion.	**Azure:** Azure Data Explorer **GCP:** Bigtable / BigQuery **Neutral:** InfluxDB / TimescaleDB	1-3 months
EventBridge Integration AI agent help: high $1.7k/mo in scope	Proprietary event bus whose value is AWS-native event sources, rules, Pipes and the Scheduler. Event patterns and targets are AWS ARNs.	Re-implement routing on Event Grid / Eventarc or Kafka; replace AWS-native sources with the target's audit/event feeds; move schedules to the target scheduler.	**Azure:** Event Grid **GCP:** Eventarc / Cloud Scheduler **Neutral:** Kafka / Argo Events	1-2 months
Amplify DevOps AI agent help: high $1.3k/mo in scope	Opinionated framework that generates Cognito, AppSync, DynamoDB and Lambda backends; front-end hosting is portable, the generated backend is not.	Host front-ends on Static Web Apps / Firebase Hosting / Vercel; re-platform the backend along with Cognito, AppSync and DynamoDB.	**Azure:** Static Web Apps **GCP:** Firebase Hosting **Neutral:** Vercel / Netlify	1-3 months
Bedrock AgentCore AI/ML AI agent help: medium $1.2k/mo in scope	Agent runtime, memory, gateway, identity and observability services with AWS-specific APIs. Agent code written with open frameworks is portable; the hosting, memory stores and tool gateway are not.	Keep agent logic in an open framework (LangGraph, Strands, Claude Agent SDK); re-host on Azure AI Foundry Agent Service, Vertex AI Agent Engine or containers; migrate memory to a vector/relational store you control.	**Azure:** Azure AI Foundry Agent Service **GCP:** Vertex AI Agent Engine **Neutral:** Containers + open agent frameworks	1-2 months
Lake Formation Analytics AI agent help: medium $640/mo in scope	Proprietary fine-grained permission layer over the Glue catalog; access policies must be re-modelled.	Re-implement as Unity Catalog, Microsoft Purview / OneLake security, or Dataplex policies.	**Azure:** Microsoft Purview / OneLake security **GCP:** Dataplex / BigLake **Neutral:** Unity Catalog / Apache Ranger	1-2 months
Data Pipeline Analytics AI agent help: high $327/mo in scope	Legacy proprietary orchestration (in maintenance mode).	Rewrite as Airflow DAGs or Data Factory pipelines.	**Azure:** Data Factory **GCP:** Cloud Composer **Neutral:** Apache Airflow	1-2 months

Tier 3 · Moderate 10 services

Service	Why this tier	Recommended migration path	Equivalents	Typical duration
SageMaker AI/ML AI agent help: medium $23k/mo in scope	Models (PyTorch, XGBoost, Hugging Face) and containers are portable, but pipelines, feature store, endpoints and Studio domains are SageMaker-specific.	Move training/inference to Azure ML / Vertex AI or Kubeflow/MLflow; rebuild pipelines and feature store.	**Azure:** Azure Machine Learning **GCP:** Vertex AI **Neutral:** Kubeflow / MLflow / KServe	1-3 months
API Gateway Integration AI agent help: high $11k/mo in scope	OpenAPI definitions export, but Lambda/VTL integrations, authorizers, usage plans and WebSocket routes are AWS-specific.	Export OpenAPI, re-implement on Azure API Management, Apigee or Kong; rewrite authorizers and mapping templates.	**Azure:** API Management **GCP:** Apigee / API Gateway **Neutral:** Kong / Envoy Gateway	3-8 weeks
Bedrock AI/ML AI agent help: medium $7.6k/mo in scope	Foundation models (Anthropic Claude, Llama, Mistral) are available on other clouds and direct from providers, but Bedrock Agents, Knowledge Bases, Guardrails and Flows are proprietary orchestration.	Call the same models via Vertex AI / Azure AI Foundry or the provider's own API; re-implement agents and RAG pipelines with a portable framework; re-run evaluations.	**Azure:** Azure AI Foundry **GCP:** Vertex AI **Neutral:** Provider APIs (Anthropic	3-8 weeks
Neptune Database AI agent help: low $6.8k/mo in scope	Graph database using open query languages (Gremlin, openCypher, SPARQL), but loaders, IAM auth and Neptune-specific extensions need rework.	Move to Neo4j, Cosmos DB for Gremlin or JanusGraph; re-validate query behaviour.	**Azure:** Cosmos DB for Apache Gremlin **GCP:** Neo4j AuraDB / Spanner Graph **Neutral:** Neo4j / JanusGraph	1-3 months
Athena Analytics AI agent help: high $6.6k/mo in scope	Trino/Presto-based SQL, so queries are largely portable, but it is bound to the Glue Data Catalog, S3 locations and Athena workgroups.	Move tables to Iceberg/Delta, query with Trino, BigQuery or Databricks SQL.	**Azure:** Fabric / Synapse serverless SQL **GCP:** BigQuery **Neutral:** Trino / Starburst	3-8 weeks
SQS Integration AI agent help: high $3.6k/mo in scope	Simple semantics, but a proprietary API; every producer and consumer needs a code change.	Move to Service Bus / Pub/Sub / RabbitMQ behind a small messaging abstraction; drain queues at cutover.	**Azure:** Service Bus / Storage Queues **GCP:** Pub/Sub / Cloud Tasks **Neutral:** RabbitMQ / NATS	2-6 weeks

Service	Why this tier	Recommended migration path	Equivalents	Typical duration
Transcribe AI/ML AI agent help: medium $3.4k/mo in scope	Proprietary speech API; custom vocabularies need recreating.	Swap to Azure AI Speech / Speech-to-Text or Whisper.	**Azure:** Azure AI Speech **GCP:** Speech-to-Text **Neutral:** Whisper	2-4 weeks
Textract AI/ML AI agent help: medium $3.0k/mo in scope	Proprietary document-extraction API and output schema.	Move to Azure Document Intelligence / Document AI; re-map output parsing and re-validate accuracy.	**Azure:** Document Intelligence **GCP:** Document AI **Neutral:** Open-source OCR + LLM extraction	1-2 months
Firehose Analytics AI agent help: medium $2.6k/mo in scope	Managed delivery pipeline with AWS-only sources and destinations.	Replace with Kafka Connect, Event Hubs Capture, or Dataflow templates.	**Azure:** Event Hubs Capture / Stream Analytics **GCP:** Dataflow / Pub/Sub subscriptions **Neutral:** Kafka Connect / Vector	3-6 weeks
SNS Integration AI agent help: high $1.8k/mo in scope	Proprietary pub/sub and SMS/mobile push API, usually fanned out to SQS and Lambda.	Move topics to Event Grid / Service Bus topics / Pub/Sub; move SMS to a CPaaS provider (Twilio, Azure Communication Services).	**Azure:** Event Grid / Service Bus topics **GCP:** Pub/Sub **Neutral:** NATS / Kafka + Twilio	2-6 weeks

Tier 2 · Low 18 services

Service	Why this tier	Recommended migration path	Equivalents	Typical duration
RDS / Aurora Database AI agent help: low $107k/mo in scope	PostgreSQL, MySQL, MariaDB, SQL Server and Oracle engines, plus Aurora (PostgreSQL/MySQL compatible). Schemas and SQL move; Aurora-only features (Global Database, Babelfish, Aurora Serverless scaling, IAM DB auth) and Oracle/SQL Server licensing need attention.	Migrate with native replication or DMS-equivalents (Azure DMS, Database Migration Service) into Azure Database / Cloud SQL / AlloyDB. Cut over per database with a short read-only window.	**Azure:** Azure Database for PostgreSQL/MySQL / Azure SQL **GCP:** Cloud SQL / AlloyDB **Neutral:** Self-managed PostgreSQL (CloudNativePG)	2-8 weeks per database
S3 Storage AI agent help: low $55k/mo in scope	The S3 API is a de-facto standard implemented by most object stores, but bucket policies, event notifications, lifecycle rules, Object Lock and presigned URL patterns need translating, and large data sets incur egress cost and transfer time.	Copy with rclone / Storage Transfer Service / AzCopy; use S3-compatible endpoints where code depends on the S3 SDK; recreate lifecycle, retention and event rules.	**Azure:** Blob Storage **GCP:** Cloud Storage **Neutral:** MinIO / Ceph RGW	1-8 weeks (data-volume bound)
ECS Compute AI agent help: high $14k/mo in scope	Containers are portable but ECS task definitions, service discovery and deployment wiring are AWS-specific and must be rewritten.	Convert task definitions to Kubernetes manifests or Azure Container Apps / Cloud Run definitions. Images move unchanged.	**Azure:** Container Apps / AKS **GCP:** Cloud Run / GKE **Neutral:** Kubernetes / Nomad	2-6 weeks per service group
Direct Connect Networking AI agent help: none $11k/mo in scope	Physical circuits; lead times for new carrier connections dominate.	Order ExpressRoute / Cloud Interconnect early (8-16 week lead time is common).	**Azure:** ExpressRoute **GCP:** Cloud Interconnect **Neutral:** Carrier MPLS	2-4 months (carrier lead time)
EMR Analytics AI agent help: medium $11k/mo in scope	Open-source Spark, Hive, Presto, Flink. EMRFS, bootstrap actions and step orchestration need adjusting.	Run the same jobs on Databricks, Dataproc or HDInsight; swap s3:// paths and IAM roles.	**Azure:** Databricks / HDInsight **GCP:** Dataproc **Neutral:** Spark on Kubernetes	1-2 months
Marketplace (3rd party) Third-party software AI agent help: none $9.2k/mo in scope	Third-party software bought through AWS billing. The product usually exists on other clouds or direct from the vendor; contracts and private offers need re-papering.	Re-contract directly or via the target cloud's marketplace; align with contract renewal dates.	**Azure:** Azure Marketplace **GCP:** Google Cloud Marketplace **Neutral:** Direct vendor contract	Contract-bound
CloudWatch Management AI agent help: high $8.4k/mo in scope	Metrics, logs, alarms and dashboards are re-creatable, but Embedded Metric Format, Logs Insights queries and alarm-driven automation are AWS-specific.	Standardise on OpenTelemetry and ship to Azure Monitor / Cloud Monitoring / Grafana stack; rebuild alarms and dashboards from code.	**Azure:** Azure Monitor **GCP:** Cloud Monitoring / Logging **Neutral:** Grafana / Prometheus / Loki / Datadog	2-4 weeks
CloudFront Networking AI agent help: medium $7.4k/mo in scope	A CDN like any other, but Lambda@Edge / CloudFront Functions and signed-URL schemes are proprietary.	Move to Front Door, Cloud CDN, Cloudflare or Akamai; rewrite edge functions.	**Azure:** Front Door **GCP:** Cloud CDN **Neutral:** Cloudflare / Akamai / Fastly	1-3 weeks
Network Firewall Networking AI agent help: medium $5.3k/mo in scope	Suricata-compatible rules; policy structure is AWS-specific.	Translate rules to Azure Firewall / Cloud NGFW or a vendor NGFW.	**Azure:** Azure Firewall **GCP:** Cloud NGFW **Neutral:** Palo Alto / Fortinet	2-4 weeks
Managed Flink Analytics AI agent help: high $3.6k/mo in scope	Apache Flink applications are portable; legacy Kinesis SQL applications are not.	Run Flink jobs on Confluent, Ververica, or Flink on Kubernetes; rewrite any legacy SQL apps.	**Azure:** Flink on HDInsight on AKS / Stream Analytics **GCP:** Dataflow / Flink on GKE **Neutral:** Apache Flink	1-2 months
Backup Storage AI agent help: low $3.4k/mo in scope	Backup policies and vaults are AWS-specific; the protected data moves with each workload. Retained backups must be kept (or restored and re-protected) to meet retention obligations.	Recreate policies in Azure Backup / Backup and DR; keep the AWS vault read-only until retention expires.	**Azure:** Azure Backup **GCP:** Backup and DR **Neutral:** Veeam / Commvault	Retention-bound tail

Service	Why this tier	Recommended migration path	Equivalents	Typical duration
KMS Security AI agent help: none $2.4k/mo in scope	Keys cannot be exported from KMS. Data encrypted under KMS must be decrypted and re-encrypted on the target; envelope-encryption code calls the KMS API.	Create keys in Key Vault / Cloud KMS (or an external HSM / BYOK), re-encrypt data during migration, swap the SDK calls.	**Azure:** Key Vault / Managed HSM **GCP:** Cloud KMS / Cloud HSM **Neutral:** HashiCorp Vault / Thales HSM	Moves with the workload
Directory Service Security AI agent help: medium $2.2k/mo in scope	Managed Microsoft AD - standard Active Directory.	Extend trusts to Entra Domain Services / Managed AD on the target and migrate members.	**Azure:** Entra Domain Services **GCP:** Managed Service for Microsoft AD **Neutral:** Active Directory	1-2 months
ACM / Private CA Security AI agent help: low $1.6k/mo in scope	Public certificates are free to reissue elsewhere; a private CA hierarchy must be re-established and trusted by every client.	Reissue public certs on the target; plan a private CA trust rollout (or adopt a cloud-neutral CA such as Vault PKI).	**Azure:** Key Vault certificates **GCP:** Certificate Authority Service **Neutral:** Vault PKI / EJBCA	2-6 weeks
WAF Security AI agent help: medium $1.5k/mo in scope	Rule concepts are common; rule syntax and managed rule groups are AWS-specific.	Translate rules to Azure WAF / Cloud Armor / Cloudflare WAF; re-tune in detection mode before enforcing.	**Azure:** Azure WAF **GCP:** Cloud Armor **Neutral:** Cloudflare / Akamai WAF	2-4 weeks
Global Accelerator Networking AI agent help: medium $766/mo in scope	Anycast IP front door; static IPs cannot be moved.	Replace with Front Door / global load balancer; plan for client IP allow-list changes.	**Azure:** Front Door / cross-region Load Balancer **GCP:** Global external Load Balancer **Neutral:** Cloudflare Spectrum	2-4 weeks
Secrets Manager Security AI agent help: medium $647/mo in scope	Simple secret store, but rotation Lambdas and SDK calls are AWS-specific.	Move secrets to Key Vault / Secret Manager / Vault; rewrite rotation; prefer External Secrets Operator for Kubernetes workloads.	**Azure:** Key Vault **GCP:** Secret Manager **Neutral:** HashiCorp Vault	1-2 weeks
Batch Compute AI agent help: high $0/mo in scope	Container jobs on EC2/Fargate; job definitions and queues are AWS-specific.	Move to Azure Batch, Google Batch, or Kubernetes Jobs / Argo Workflows.	**Azure:** Azure Batch **GCP:** Google Batch **Neutral:** Kubernetes Jobs / Slurm	2-4 weeks

Tier 1 · Portable 12 services

Service	Why this tier	Recommended migration path	Equivalents	Typical duration
EC2 Compute AI agent help: low $188k/mo in scope	Virtual machines running standard operating systems. Images can be exported or rebuilt from configuration management.	Rehost VMs (Azure Migrate, Google Migrate to VMs, or rebuild from Packer/Ansible images). Re-baseline instance sizes on the target.	**Azure:** Virtual Machines **GCP:** Compute Engine **Neutral:** VMware / OpenStack / bare metal	1-4 weeks per application
EC2 Other (EBS, NAT, transfer) Compute AI agent help: low $43k/mo in scope	Block storage, snapshots, NAT gateways, elastic IPs and data transfer. Commodity primitives on every cloud.	Recreate volumes and network egress on the target; move data with snapshots or block replication. Budget for one-off egress charges.	**Azure:** Managed Disks / NAT Gateway **GCP:** Persistent Disk / Cloud NAT **Neutral:** SAN / Ceph	Moves with the workload
EKS Compute AI agent help: medium $14k/mo in scope	Upstream Kubernetes. Manifests and Helm charts run on any conformant cluster; AWS-specific add-ons (IRSA, ALB controller, EBS CSI) are small and swappable.	Stand up AKS/GKE/on-prem Kubernetes, swap cloud add-ons (identity, ingress, storage classes), redeploy via GitOps.	**Azure:** AKS **GCP:** GKE **Neutral:** OpenShift / Rancher / upstream Kubernetes	2-6 weeks per cluster
VPC Networking AI agent help: medium $12k/mo in scope	Standard network constructs (subnets, routing, VPN, IPAM); only the console and APIs differ.	Design the target landing-zone network; reuse the IP plan where possible.	**Azure:** Virtual Network **GCP:** VPC **Neutral:** Data-centre networking	Part of landing zone
MSK (Kafka) Analytics AI agent help: low $7.9k/mo in scope	Standard Apache Kafka.	Mirror topics with MirrorMaker 2 / Cluster Linking and cut consumers over.	**Azure:** Event Hubs (Kafka) / Confluent **GCP:** Managed Service for Apache Kafka **Neutral:** Apache Kafka / Redpanda	2-6 weeks
FSx Storage AI agent help: none $5.7k/mo in scope	Managed Windows File Server, Lustre, ONTAP or OpenZFS - all standard file systems.	Replicate to the equivalent managed file system (ONTAP SnapMirror where applicable).	**Azure:** Azure Files / NetApp Files **GCP:** Filestore / NetApp Volumes **Neutral:** NetApp / Windows File Server	2-4 weeks
OpenSearch Analytics AI agent help: low $4.7k/mo in scope	Open-source OpenSearch / Elasticsearch APIs (Serverless collections have a reduced API).	Snapshot/restore or reindex into Elastic Cloud, OpenSearch on Kubernetes or Azure AI Search where suitable.	**Azure:** Elastic Cloud on Azure **GCP:** Elastic Cloud on GCP **Neutral:** OpenSearch	2-4 weeks
ELB Networking AI agent help: medium $4.2k/mo in scope	Standard L4/L7 load balancing.	Recreate listeners and rules on Application Gateway / Cloud Load Balancing / Envoy.	**Azure:** Load Balancer / Application Gateway **GCP:** Cloud Load Balancing **Neutral:** HAProxy / F5 / Envoy	Moves with the workload
ElastiCache Database AI agent help: low $4.2k/mo in scope	Managed Redis / Valkey / Memcached - open engines.	Provision the target cache and warm it; caches rarely need data migration.	**Azure:** Azure Cache for Redis **GCP:** Memorystore **Neutral:** Redis / Valkey	Days

Service	Why this tier	Recommended migration path	Equivalents	Typical duration
MWAA (Airflow) Analytics AI agent help: medium $2.0k/mo in scope	Apache Airflow; DAGs are portable apart from AWS operators.	Move DAGs to Cloud Composer, Astronomer or Airflow on Kubernetes; swap AWS operators.	**Azure:** Data Factory Workflow Orchestration Manager / Astronomer **GCP:** Cloud Composer **Neutral:** Apache Airflow	2-4 weeks
Amazon MQ Integration AI agent help: low $1.4k/mo in scope	Managed ActiveMQ / RabbitMQ - standard protocols (AMQP, MQTT, STOMP, JMS).	Shovel/federate queues to the target broker and cut clients over.	**Azure:** Service Bus (AMQP) / RabbitMQ on AKS **GCP:** RabbitMQ on GKE **Neutral:** RabbitMQ / ActiveMQ	1-3 weeks
Route 53 Networking AI agent help: medium $991/mo in scope	DNS zones export as standard zone files. Alias records, health-check routing and Resolver rules need translating.	Export zones, recreate in Azure DNS / Cloud DNS / NS1, then lower TTLs and switch NS records.	**Azure:** Azure DNS **GCP:** Cloud DNS **Neutral:** NS1 / Cloudflare DNS	1-2 weeks

Unclassified services in this estate

Scored as Moderate (tier 3) until added to the catalog.

- AWS Transform: $495/month

B Appendix: methodology

Data sources

Spend comes from AWS Cost Explorer in the management (payer) account: amortised cost per account per service per month, excluding tax, credits, refunds and support. Marketplace purchases are grouped as one third-party line. Accounts and OU structure come from AWS Organizations. The optional deep scan counts resources (Lambda functions, DynamoDB tables, state machines, stacks and more) in each member account through a read-only role. All API calls are read-only.

Scoring

Each service's tier maps to a score: Portable 10, Low 30, Moderate 50, High 75, Critical 95. An account's **inherent risk** is the spend-weighted average of its services' scores.

The report gives **two OU scores**. **By spend** weights each account by its share of in-scope spend, so a small account full of proprietary services scores high on its own yet moves the score very little. **By number of accounts** counts every account with spend equally, so one large account (for example a contact centre on Connect) cannot dominate it. A gap between the two shows whether risk is concentrated in a few big accounts or spread across many. An account's **contribution** is its inherent risk × its weight under the weighting chosen in the bar at the top (by spend unless changed).

Ratings: Low < 25 ≤ Moderate < 40 ≤ Elevated < 55 ≤ High < 70 ≤ Critical.

Estimate

The cost and time figures depend on the delivery questions under Exit model. Until they are answered, the report marks the figures provisional and keeps them out of the bottom line.

Effort is built bottom-up: each material service in each account contributes the catalog's engineer-weeks for one workload, scaled by the square root of its spend relative to US$1,000/month (and by its resource count where the deep scan ran). Account overhead, a target landing zone, programme management and contingency are added, and the total is adjusted for delivery maturity from the questionnaire. Duration is the longer of effort ÷ available capacity and the critical path of the hardest tier present.

Limitations

- Spend is a proxy for size; a cheap service can still be hard to move (the deep scan's resource counts partly correct for this).
- RDS spend covers all engines; Aurora-specific features and commercial engines (Oracle, SQL Server) appear in the deep scan's engine breakdown.
- Contractual exit costs (EDP commitments, Savings Plans, Reserved Instances) and target-cloud run costs are outside the estimate; commitments still owed after the exit are shown separately, as written off, under Timing.
- The estimate is an order-of-magnitude planning figure for an executive decision, not a delivery quote.